

# **Introduction to High Performance Scientific Computing**

**Autumn, 2017**

**Lecture 1**

# Instructor

---

**Prasun Ray**

**Teaching Fellow**

**Department of Mathematics**

**[p.ray@imperial.ac.uk](mailto:p.ray@imperial.ac.uk)**

**Huxley 6M20**

**Office hours: Mondays 5-6pm, MLC**

**Thursdays 4-5pm, MLC**

**(First office hour on Monday, 9/10)**

# Weekly schedule

---

- **Lectures:**

**Monday, 11-12, Huxley 340**

**Thursday, 11-12, Huxley 340**

- **Labs:**

**Tuesday, 5-6pm, MLC (Huxley 414)**

**or**

**Wednesday, 10-11am, Huxley 340**

- **Only need to attend *one* lab session**

**CDT-only lab: 11am-12pm, RSM 3.38**

- **Wednesday 10am lab requires laptop with necessary software installed (more on this later)**

# Syllabus

---

**Lectures 1-2: Unix basics, version control with git/bitbucket**

**Lectures 3-6: Programming and scientific computing with Python**

**Lectures 7-10: Modular programming with Fortran, libraries, makefiles, coupling Fortran+Python**

# Syllabus

---

**Lectures 11-14: Introduction to parallel computing and OpenMP**

**Lectures 15-16: Distributed memory computing with MPI, parallel libraries**

**Lectures 17-20: Basic computer architecture, cloud computing, cluster computing with Python and Spark**

# Assessment

---

## **3 Programming assignments**

HW1: Assigned 23/10, due 2/11 (**20%**)

HW2: Assigned 6/11, due 16/11 (**20%**)

HW3: Assigned 20/11, due 29/11 (**15%**)

## **1 Programming Project (45%)**

Assigned 30/11, due 15/12

**Submitting HW1 commits you to the course**

**CDT students:** Will be contacted about assessment separately

# Online material

---

- Main resource is course webpage:

<http://imperialhpsc.bitbucket.io/>


- Slides will be available before every lecture

All course material will be available on course bitbucket page  
(more on this later):

<https://bitbucket.org/ImperialHPSC/m3c2017>

# Moore's law

## Microprocessor Transistor Counts 1971-2011 & Moore's Law




- Number of transistors on chip doubles every two years
- From ~2005, multicore processors (why?)
- 2 core, then 4, 6, 8, 16
- Intel Xeon Phi: 60+ cores!



# Algorithms and hardware


Speed-Up  
Factor

Derived from Computational Methods



Speed-Up  
Factor

Derived from Supercomputer Hardware



# High-end HPC

RANK	SITE	SYSTEM	CORES	RMAX (TFLOP/S)	RPEAK (TFLOP/S)	POWER (KW)
1	National Super Computer Center in Guangzhou China	<b>Tianhe-2 (MilkyWay-2)</b> - TH-IVB-FEP Cluster, Intel Xeon E5-2692 12C 2.200GHz, TH Express-2, Intel Xeon Phi 31S1P NUDT	3,120,000	33,862.7	54,902.4	17,808
2	DOE/SC/Oak Ridge National Laboratory United States	<b>Titan</b> - Cray XK7 , Opteron 6274 16C 2.200GHz, Cray Gemini interconnect, NVIDIA K20x Cray Inc.	560,640	17,590.0	27,112.5	8,209
3	DOE/NNSA/LLNL United States	<b>Sequoia</b> - BlueGene/Q, Power BQC 16C 1.60 GHz, Custom IBM	1,572,864	17,173.2	20,132.7	7,890
4	RIKEN Advanced Institute for Computational Science (AICS) Japan	K computer, SPARC64 VIIIfx 2.0GHz, Tofu interconnect Fujitsu	705,024	10,510.0	11,280.4	12,660

**Historically: cluster computing limited to national labs, research universities**

**But now...**



# Cluster computing is mainstream

---



Google Cloud Platform



# Course objective

---

- Cluster computing is not free!
- Important to:
  - choose right tools
  - use them effectively

*This course provides foundation for “intelligent, informed” computing.*

# Software tools

---

Useful to classify tools as *scientific* or *general purpose*

# Software tools

---

Useful to classify tools as *scientific* or *general purpose*

Examples:

Scientific	General purpose
Matlab	Python
Fortran	C++
R	Java

# Software tools

---

Languages are *compiled* or *interpreted*



# Software tools

---

Languages are *compiled* or *interpreted*

Compiled	Interpreted
Fortran	Python
C++	Matlab
Java	R

# Software tools

---

**This course:**

**Python: interpreted, general purpose**

**Fortran: compiled, scientific**

# Operating systems

---

**Most HPC and scientific computing requires Unix (or Unix-like terminals)**

**Linux and Mac OS are built on Unix (and have terminal apps)**

- **Fairly straightforward to install course software**

**Windows:**


- **Not well-suited for HPC**
- **Can get Unix terminal with cygwin**
- **For this course: Should install Linux virtual machine (VM) and install software within the VM**
- **MLC computers have Linux VMs installed (go try them out!)**

**Instructions for installing course software available online:**

**<http://imperialhpsc.bitbucket.io/>**

# Unix terminal

## Terminal on a mac:



# 12 Unix commands

---

## Navigation:

*pwd*: print working directory (where am I?)

*ls*: list of directory contents (what is here?)

*cd*: change directory (let's go somewhere else)

```
$ pwd
/Users/prasun/Documents/repos/m3c2017
$
$ ls
Readme.md lectures
$
$ cd lectures
$
$ ls
lecture1
```

# 12 Unix commands

---

## Manipulate files and directories:

*cp*: Make copy of a file

*mv*: Move or rename a file

*rm*: Remove a file

*rm -r* : Remove directory and  
all of its contents (dangerous!)

```
$ ls
Readme.md lectures
$
$ cp Readme.md Readme.md_copy
$
$ ls
Readme.md Readme.md_copy
        lectures
$
$ mv Readme.md_copy
Readme.md_copy2
$
$ ls
Readme.md Readme.md_copy2
        lectures
$
$ rm Readme.md_copy2
$
$ ls
Readme.md lectures
```

# 12 Unix commands

---

## Info about contents of file:

*cat*: List contents of file

*head -n*: List first n lines

*tail -n*: list last n lines

*grep*: search within file for a string

```
$ cat example.txt
This is an example text file.
This is line 2.
This is line 3.
This is the last line.
$
$ head -1 example.txt
This is an example text file.
$
$ tail -2 example.txt
This is line 3.
This is the last line.
$
$ grep last example.txt
This is the last line.
```

# 12 Unix commands

---

**Getting help:**

*man*: manual page for a command

**Try** *man ls*. What does *ls -l* do? *ls -a*?

**What if you don't know name of command?**

[https://en.wikipedia.org/wiki/List\\_of\\_Unix\\_commands](https://en.wikipedia.org/wiki/List_of_Unix_commands)

**or google.**



# 12 Unix commands

---

The 12 commands:

1. `pwd`
2. `ls`
3. `cd`
4. `cp`
5. `mv`
6. `rm`
7. `rm -r`
8. `cat`
9. `head -n`
10. `tail -n`
11. `grep`
12. `man`

This is “basic” Unix. Can do much more!

# A little more Unix

---

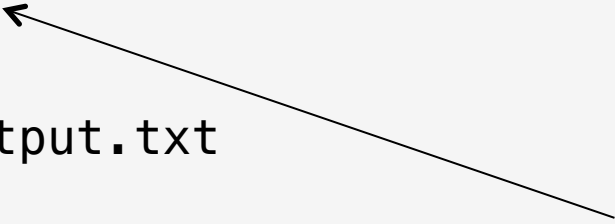
Instead of outputting to screen, can output to file using “>”

```
$ ls
example.txt  lecture1

$ grep last example.txt > output.txt

$ ls
example.txt  lecture1  output.txt

$ cat output.txt
This is the last line.
```



**Lines in example.txt containing “last”  
are written to output.txt**

# A little more Unix

---

**Command can be executed sequentially (they can be “piped”) using “|”**

```
$ head -2 example.txt | grep line > output.txt  
$  
$ cat output.txt  
This is line 2.
```



**First two lines in example.txt are searched for the string “line” with results being written to output.txt**

# An example

---

You run optimization software that gives output that looks like:

```
INPUT:endgeom
INPUT:azimuthal 9 0.1
INPUT:polar 5
INPUT:begin
```

```
k-cactus is 1.402458
```

```
TIMING: Module: cpu      10.03 wall      10.04 Overall: cpu      29.00 wall      29.29
```

```
=====
=====
```

```
INPUT:EDIT 4
```

```
CALLING EDIT(INTERFACE_NO= 4)
```

```
INPUT:begin
```

```
*****
```

```
INTERFACE  4 EIGENVALUE 1.402458 OVERALL MWd/t  0.0000E+00 BURNUP TIME  0.0000E+00
DAYS
```

```
=====
=====
```

```
RUN SET  1
Imperial College
London
```

# An example

**We only care about the “k-cactus” values which appear several times.  
How do we extract them?**

```
INPUT:endgeom
INPUT:azimuthal 9 0.1
INPUT:polar 5
INPUT:begin
```

```
k-cactus is 1.402458
```

```
TIMING: Module: cpu      10.03 wall      10.04 Overall: cpu      29.00 wall      29.29
```

```
=====
=====
```

```
INPUT:EDIT 4
```

```
CALLING EDIT(INTERFACE_NO= 4)
```

```
INPUT:begin
```

```
*****
```

```
INTERFACE  4 EIGENVALUE 1.402458 OVERALL MWd/t  0.0000E+00 BURNUP TIME  0.0000E+00
DAYS
```

```
=====
=====
```

```
RUN SET  1
Imperial College
London
```

# An example

---

## Using grep:

```
$ grep cactus datafile.out
k-cactus is      1.402458
k-cactus is      1.386050
k-cactus is      1.377296
k-cactus is      1.352324
k-cactus is      1.328779
```

**But what if we only want the numbers?**

# An example

---

## Using grep:

```
$ grep cactus datafile.out
k-cactus is      1.402458
k-cactus is      1.386050
k-cactus is      1.377296
k-cactus is      1.352324
k-cactus is      1.328779
```

## But what if we only want the numbers?

Use “cut”:

```
$ grep cactus datafile.out | cut -d s -f 3
1.402458
1.386050
1.377296
1.352324
1.328779
```

**Questions:** How do we store these numbers in a file? How do we find out what the flags after “cut” are doing?

# What next?

---

- **If you have your own laptop/desktop: start installing course software – see webpage for instructions**
- **If you don't, try out the virtual machines and these Unix commands on the MLC computers -- look for the “oel” icon on the desktop**
- **Start working through introductory Python videos and exercises**
  - **1<sup>st</sup> video should be online Friday morning**