

Introduction to High Performance Scientific Computing

Autumn, 2017

Lecture 1

Imperial College
London

5 October, 2017

Instructor

Prasun Ray
 Teaching Fellow
 Department of Mathematics
p.ray@imperial.ac.uk
 Huxley 6M20
 Office hours: Mondays 5-6pm, MLC
 Thursdays 4-5pm, MLC
 (First office hour on Monday, 9/10)

Imperial College
London

Weekly schedule

- Lectures:**
 Monday, 11-12, Huxley 340
 Thursday, 11-12, Huxley 340
- Labs:**
 Tuesday, 5-6pm, MLC (Huxley 414)
 or
 Wednesday, 10-11am, Huxley 340
- Only need to attend *one* lab session**
 CDT-only lab: 11am-12pm, RSM 3.38
- Wednesday 10am lab requires laptop with necessary software installed (more on this later)**

Imperial College
London

Syllabus

Lectures 1-2: Unix basics, version control with git/bitbucket

Lectures 3-6: Programming and scientific computing with Python

Lectures 7-10: Modular programming with Fortran, libraries, makefiles, coupling Fortran+Python

Imperial College
London

Syllabus

Lectures 11-14: Introduction to parallel computing and OpenMP

Lectures 15-16: Distributed memory computing with MPI, parallel libraries

Lectures 17-20: Basic computer architecture, cloud computing, cluster computing with Python and Spark

Imperial College
London

Assessment

3 Programming assignments
 HW1: Assigned 23/10, due 2/11 (20%)
 HW2: Assigned 6/11, due 16/11 (20%)
 HW3: Assigned 20/11, due 29/11 (15%)

1 Programming Project (45%)
 Assigned 30/11, due 15/12

Submitting HW1 commits you to the course

CDT students: Will be contacted about assessment separately

Imperial College
London

Online material

- Main resource is course webpage:

<http://imperialhpsc.bitbucket.io/>

- Slides will be available before every lecture


All course material will be available on course bitbucket page (more on this later):

<https://bitbucket.org/ImperialHPSC/m3c2017>

Imperial College
London

Moore's law


Microprocessor Transistor Counts 1971-2011 & Moore's Law



Imperial College
London

"Transistor Count and Moore's Law - 2011" by Wigimom - Own work. Licensed under CC BY-SA 3.0 via Wikimedia Commons - http://commons.wikimedia.org/wiki/File:Transistor_Count_and_Moore%27s_Law_-_2011.png#/media/File:Transistor_Count_and_Moore%27s_Law_-_2011.png

Algorithms and hardware



Imperial College
London

High-end HPC

RANK	SITE	SYSTEM	CORES	RMAX (TFLOP/S)	RPEAK (TFLOP/S)	POWER (KW)
1	National Super Computer Center in Guangzhou China	Tianhe-2 (Milky Way-2) - TH-WF-FEP Cluster, Intel Xeon E5-2692 12C 2.20GHz; TH Express-2, Intel Xeon Phi 3151P NUDT	3,120,000	33,862.7	54,902.4	17,808
2	DOE/SOAK Ridge National Laboratory United States	Titan - Cray XK7 , Opteron 6274 14C 2.00GHz; Cray Gemini interconnect, NVIDIA K20X Cray Inc.	560,640	17,590.0	27,112.5	8,209
3	DOE/NNSA/LLNL United States	Sequoia - BlueGene/Q, Power BGC 14C 1.60 GHz, Custom IBM	1,572,864	17,173.2	20,132.7	7,890
4	RIKEN Advanced Institute for Computational Science (AICS)	K computer, SPARC64 VIIItx 2.0GHz, Tofu interconnect Fujitsu	705,024	10,510.0	11,280.4	12,640

Historically: cluster computing limited to national labs, research universities

But now...

Imperial College
London

Cluster computing is mainstream

Big data means big computers!



Imperial College
London

Cluster computing is mainstream

Imperial College
London

Course objective

- Cluster computing is not free!
- Important to:
 - choose right tools
 - use them effectively

This course provides foundation for "intelligent, informed" computing.

Imperial College
London

Software tools

Useful to classify tools as *scientific* or *general purpose*

Imperial College
London

Software tools

Useful to classify tools as *scientific* or *general purpose*

Examples:

Scientific	General purpose
Matlab	Python
Fortran	C++
R	Java

Imperial College
London

Software tools

Languages are *compiled* or *interpreted*

Imperial College
London

Software tools

Languages are *compiled* or *interpreted*

Compiled	Interpreted
Fortran	Python
C++	Matlab
Java	R

Imperial College
London

Software tools

This course:

Python: interpreted, general purpose

Fortran: compiled, scientific

Imperial College
London

Operating systems

Most HPC and scientific computing requires Unix (or Unix-like terminals)

Linux and Mac OS are built on Unix (and have terminal apps)

- Fairly straightforward to install course software

Windows:


- Not well-suited for HPC
- Can get Unix terminal with cygwin
- For this course: Should install Linux virtual machine (VM) and install software within the VM
- MLC computers have Linux VMs installed (go try them out!)

Instructions for installing course software available online:
<http://imperialhpsc.bitbucket.io/>

Imperial College
London

Unix terminal

Terminal on a mac:



Imperial College
London

12 Unix commands

Navigation:

pwd: print working directory (where am I?)
ls: list of directory contents (what is here?)
cd: change directory (let's go somewhere else)

```
$ pwd
/Users/prasun/Documents/repos/m3c2017
$
$ ls
Readme.md lectures
$
$ cd lectures
$
$ ls
lecture1
```

Imperial College
London

12 Unix commands

Manipulate files and directories:

cp: Make copy of a file

mv: Move or rename a file

rm: Remove a file

rm -r: Remove directory and all of its contents (dangerous!)

```
$ ls
Readme.md lectures
$
$ cp Readme.md Readme.md_copy
$
$ ls
Readme.md Readme.md_copy
lectures
$
$ mv Readme.md_copy
Readme.md_copy2
$
$ ls
Readme.md Readme.md_copy2
lectures
$
$ rm Readme.md_copy2
$
$ ls
Readme.md lectures
```

Imperial College
London

12 Unix commands

Info about contents of file:

cat: List contents of file

head -n: List first n lines

tail -n: list last n lines

grep: search within file for a string

```
$ cat example.txt
This is an example text file.
This is line 2.
This is line 3.
This is the last line.
$
$ head -1 example.txt
This is an example text file.
$
$ tail -2 example.txt
This is line 3.
This is the last line.
$
$ grep last example.txt
This is the last line.
```

Imperial College
London

12 Unix commands

Getting help:

man: manual page for a command

Try **man ls**. What does **ls -l** do? **ls -a**?

What if you don't know name of command?

https://en.wikipedia.org/wiki/List_of_Unix_commands

or google.

Imperial College
London

12 Unix commands

The 12 commands:

1. pwd
2. ls
3. cd
4. cp
5. mv
6. rm
7. rm -r
8. cat
9. head -n
10. tail -n
11. grep
12. man

This is “basic” Unix. Can do much more!

Imperial College
London

A little more Unix

Instead of outputting to screen, can output to file using “>”

```
$ ls
example.txt  lecture1

$ grep last example.txt > output.txt

$ ls
example.txt  lecture1  output.txt

$ cat output.txt
This is the last line.
```

Lines in example.txt containing “last” are written to output.txt

Imperial College
London

A little more Unix

Command can be executed sequentially (they can be “piped”) using “|”

```
$ head -2 example.txt | grep line > output.txt
$
$ cat output.txt
This is line 2.
```

First two lines in example.txt are searched for the string “line” with results being written to output.txt

Imperial College
London

An example

You run optimization software that gives output that looks like:

```
INPUT:ndgeom
INPUT:azimuthal 9 0.1
INPUT:polar 5
INPUT:begin

k-cactus is 1.402458

TIMING: Module: cpu      10.03 wall      10.04 Overall: cpu      29.00 wall      29.29
=====
INPUT:EDIT 4

CALLING EDIT(INTERFACE_NO= 4)

INPUT:begin
=====
INTERFACE  4 EIGENVALUE 1.402458 OVERALL MWd/t 0.0000E+00 BURNUP TIME 0.0000E+00
DAYS
=====
RUN SET 1
Imperial College
London
```

An example

We only care about the “k-cactus” values which appear several times.
How do we extract them?

```
INPUT:ndgeom
INPUT:azimuthal 9 0.1
INPUT:polar 5
INPUT:begin

k-cactus is 1.402458

TIMING: Module: cpu      10.03 wall      10.04 Overall: cpu      29.00 wall      29.29
=====
INPUT:EDIT 4

CALLING EDIT(INTERFACE_NO= 4)

INPUT:begin
=====
INTERFACE  4 EIGENVALUE 1.402458 OVERALL MWd/t 0.0000E+00 BURNUP TIME 0.0000E+00
DAYS
=====
RUN SET 1
Imperial College
London
```

An example

Using grep:

```
$ grep cactus datafile.out
k-cactus is 1.402458
k-cactus is 1.386050
k-cactus is 1.377296
k-cactus is 1.352324
k-cactus is 1.328779
```

But what if we only want the numbers?

Imperial College
London

An example

Using grep:

```
$ grep cactus datafile.out
k-cactus is 1.402458
k-cactus is 1.386050
k-cactus is 1.377296
k-cactus is 1.352324
k-cactus is 1.328779
```

But what if we only want the numbers?

Use "cut": `$ grep cactus datafile.out | cut -d s -f 3`

```
1.402458
1.386050
1.377296
1.352324
1.328779
```

Questions: How do we store these numbers in a file? How do we find out what the flags after "cut" are doing?

Imperial College
London

What next?

- If you have your own laptop/desktop: start installing course software -- see webpage for instructions
- If you don't, try out the virtual machines and these Unix commands on the MLC computers -- look for the "oel" icon on the desktop
- Start working through introductory Python videos and exercises
 - 1st video should be online Friday morning

Imperial College
London